


## Spatial locality

transformer, known rule (L1)



1D positions

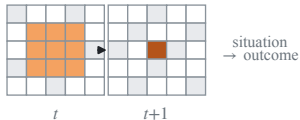


2D positions



## Temporal locality

transformer, unseen rules (L3)



own cell




+ previous-frame 3×3



## Temporal stability

same diffusion weights, known rule (L1)

settled frames feed the next



joint denoising, 8 calls



causal freezing, 8 calls



sequence-level accuracy (%), means over three seeds